


Data Wrangling with `tidyr`

Stephanie Hicks, Rafael Irizarry

The data analysis process can be thought about in four parts

1. Data cleaning
2. Data transformation
3. Data visualization
4. Modeling

where each of these steps need their own tools and software to complete.



As we have seen in class, one of the most time-consuming aspects of the data analysis process is “data wrangling”. This is also known as “data munging”, which is a trendy term for *cleaning up a messy data set*. This refers to the first two steps in the data analysis process:

1. Data cleaning (or tidying data)
2. Data transformation

It can take a long time to clean and transform messy data into a format that is useful for data visualization and modeling, but there are tools that can help turn messy data into clean data.

Defining data structures

There are many ways to define the structure of a data set. Most data frames are made up of **rows** and **columns** where the columns are almost always labeled and the rows are *sometimes* labeled.

For example, a data set could be structured in the following way:

- each row represents one company (row names are companies)
- each column represent one time point
- the stock prices are defined for each row/column pair

> stocks

	2016-01-01	2016-01-02	2016-01-03	2016-01-04	2016-01-05
Google	99.43952	99.76982	101.55871	100.07051	100.12929
Facebook	103.43013	100.92183	97.46988	98.62629	99.10868
Twitter	104.89633	101.43926	101.60309	100.44273	97.77664

Alternatively, a data set can be structured in the following way:

- each row represents one time point (but no row names)
- the first column defines the time variable and the last three columns contain the stock prices for three companies

> stocks

	time	Google	Facebook	Twitter
1	2016-01-01	99.43952	103.43013	104.89633
2	2016-01-02	99.76982	100.92183	101.43926
3	2016-01-03	101.55871	97.46988	101.60309
4	2016-01-04	100.07051	98.62629	100.44273
5	2016-01-05	100.12929	99.10868	97.77664

In both cases, the data is the same, but the structure is different. This can be *frustrating* to deal with because the meaning of the values (rows and columns) in the two data sets are different. Providing a standardized way of organizing values within a data set would alleviate a major portion of this frustration.

Defining tidy data

Now, we will introduce the concept of **tidy** data. Tidy data is a standard way of mapping the meaning of a dataset to its structure. The properties of a tidy data set are based on:

- Each column is a variable
- Each row is an observation

Working with tidy data is useful because it creates a structured way of organizing data values within a data set. This makes the data analysis process more efficient and simplifies the development of data analysis tools that work together. In this way, you can focus on the problem you are investigating, rather than the uninteresting logistics of data.

What is tidyr?

tidyr is an R package that transforms data sets to a tidy format.

There are two main functions in **tidyr**:

- **gather()** = takes multiple columns, and gathers them into key-value pairs (it makes “wide” data longer)
- **separate()** = turns a single character column into multiple columns (it makes “long” data wider)

We’ll explore what it means to go between a “wide” and “long” data format using **gather()** and **separate()** next.

How do I get tidyr?

To install **tidyr**

```
install.packages("tidyr")
```

To load `tidyr` and we'll need `dplyr`

```
library(tidyr)
library(dplyr)
```

For motivation, a tidy version of the stock data we looked at above looks like this: (we'll learn how the functions work in just a moment)

```
> stocks %>%
+   gather(company, price, Google:Twitter)
```

	time	company	price
1	2016-01-01	Google	99.43952
2	2016-01-02	Google	99.76982
3	2016-01-03	Google	101.55871
4	2016-01-04	Google	100.07051
5	2016-01-05	Google	100.12929
6	2016-01-01	Facebook	103.43013
7	2016-01-02	Facebook	100.92183
8	2016-01-03	Facebook	97.46988
9	2016-01-04	Facebook	98.62629
10	2016-01-05	Facebook	99.10868
11	2016-01-01	Twitter	104.89633
12	2016-01-02	Twitter	101.43926
13	2016-01-03	Twitter	101.60309
14	2016-01-04	Twitter	100.44273
15	2016-01-05	Twitter	97.77664

In this “tidy” data set, we have three columns representing three variables (time, company name and stock price). Every row represents contains one stock price from a particular time and for a specific company.

Pipe operator: `%>%`

We have introduced the operator: `%>%`. `dplyr` imports this operator from another package ([magrittr](#) [see help file here](#)). This operator allows you to pipe the output from one function to the input of another function. Instead of nesting functions (reading from the inside to the outside), the idea of of piping is to read the functions from left to right.

Now in this case, we pipe the `stocks` data frame to the function that will gather multiple columns into key-value pairs.

Data

2016 Iowa Presidential Caucus

We will explore [public poll data from HuffPost Pollster](#) from the 2016 Iowa [Democratic](#) and [Republican](#) Presidential Caucus.

First we will read in the data:

```
library(readr)
dem_polls = read_csv("http://elections.huffingtonpost.com/pollster/2016-iowa-presidential-democratic-caucus")
rep_polls = read_csv("http://elections.huffingtonpost.com/pollster/2016-iowa-presidential-republican-caucus")
```

Let's take a look at data

```
View(dem_polls)
View(rep_polls)

glimpse(dem_polls)
glimpse(rep_polls)
```

We see there is a lot of information in each data frame. First let's use `dplyr` to select a subset of the columns.

```
dem_polls <- dem_polls %>%
  select(Pollster, `End Date`, Clinton:Undecided)

rep_polls <- rep_polls %>%
  select(Pollster, `End Date`, Trump:Walker)
```

In the democratic and republican polling data sets, there is one column representing the polling percentages for each candidate, similar to the stock price data set with multiple columns representing different companies. To **tidy** it, we need to *gather* these columns into a two-column *key-value* pair. This is often described as transforming a *wide* data set into a *long* data set.

gather()

This function gathers multiple columns and collapses them into new *key-value* pairs. This transform data from *wide* format into a *long* format.

- The **key** is the name of the *new* column that you are creating which contains the values of the column headings that you are gathering
- The **value** is the name of the *new* column that will contain the values themselves
- The third argument defines the columns to gather

```
dem_polls %>%
  gather(key = candidate, value = percentage, Clinton:Undecided)
```

```
## Source: local data frame [648 x 4]
##
##           Pollster   End Date candidate percentage
```

```
##              (chr)      (date)      (chr)      (int)
## 1      Emerson College Polling Society 2016-01-31  Clinton      51
## 2              Quinnipiac 2016-01-31  Clinton      46
## 3  Des Moines Register/Bloomberg/Selzer 2016-01-29  Clinton      45
## 4      Gravis Marketing/One America News 2016-01-27  Clinton      53
## 5              PPP (D-Progress Iowa) 2016-01-27  Clinton      48
## 6              NBC/WSJ/Marist 2016-01-26  Clinton      48
## 7              Monmouth University 2016-01-26  Clinton      47
## 8              ARG 2016-01-24  Clinton      45
## 9              Quinnipiac 2016-01-24  Clinton      45
## 10             Iowa State/WHO-HD 2016-01-22  Clinton      47
## ..              ...              ...              ...
```

To select a range of columns by name, use the “:” (colon) operator

Assessment Using the democratic poll data, apply the `gather()` function to tidy the poll data by *excluding* the Pollster and End Date columns, rather than directly providing the column names to `gather`.

Hint: Look at the `gather()` help file on how to exclude column names.

```
## Provide your code here
```

```
dem_polls %>%
  gather(key = candidate, value = percentage, -c(Pollster, `End Date`))
```

```
## Source: local data frame [648 x 4]
```

```
##
##              Pollster  End Date candidate percentage
##              (chr)      (date)      (chr)      (int)
## 1      Emerson College Polling Society 2016-01-31  Clinton      51
## 2              Quinnipiac 2016-01-31  Clinton      46
## 3  Des Moines Register/Bloomberg/Selzer 2016-01-29  Clinton      45
## 4      Gravis Marketing/One America News 2016-01-27  Clinton      53
## 5              PPP (D-Progress Iowa) 2016-01-27  Clinton      48
## 6              NBC/WSJ/Marist 2016-01-26  Clinton      48
## 7              Monmouth University 2016-01-26  Clinton      47
## 8              ARG 2016-01-24  Clinton      45
## 9              Quinnipiac 2016-01-24  Clinton      45
## 10             Iowa State/WHO-HD 2016-01-22  Clinton      47
## ..              ...              ...              ...
```

```
## To select all the columns *except* a specific column,
## use the "-" (subtraction) operator (also known as negative indexing)
```

Assessment Using the “tidy” democratic poll data, use `dplyr` to filter for only the following candidates (Clinton, Sanders, O’Malley) and for polls only ending after May 1, 2015.

```
## Provide your code here
```

```
dem_polls %>%
  gather(key = candidate, value = percentage, Clinton:Undecided) %>%
  filter(candidate %in% c("Clinton", "Sanders", "O'Malley") &
    `End Date` >= "2015-05-01")
```

```
## Source: local data frame [195 x 4]
##
##           Pollster   End Date candidate percentage
##           (chr)      (date)      (chr)      (int)
## 1 Emerson College Polling Society 2016-01-31 Clinton      51
## 2 Quinnipiac 2016-01-31 Clinton      46
## 3 Des Moines Register/Bloomberg/Selzer 2016-01-29 Clinton      45
## 4 Gravis Marketing/One America News 2016-01-27 Clinton      53
## 5 PPP (D-Progress Iowa) 2016-01-27 Clinton      48
## 6 NBC/WSJ/Marist 2016-01-26 Clinton      48
## 7 Monmouth University 2016-01-26 Clinton      47
## 8 ARG 2016-01-24 Clinton      45
## 9 Quinnipiac 2016-01-24 Clinton      45
## 10 Iowa State/WHO-HD 2016-01-22 Clinton      47
## ..           ...           ...           ...           ...
```


Assessment (optional) Using the tidy and filtered democratic poll data set, use `ggplot2` to plot the results from each poll (percentage) for each of the candidates. Color the lines by the candidate.

```
## Provide your code here
```

```
library(ggplot2)
```

```
## Warning: package 'ggplot2' was built under R version 3.2.4
```

```
dem_polls %>%
  gather(key = candidate, value = percentage, Clinton:Undecided) %>%
  filter(candidate %in% c("Clinton", "Sanders", "O'Malley") &
    `End Date` >= "2015-05-01") %>%
  ggplot(aes(x=`End Date`, y = percentage, color = candidate)) +
  geom_line()
```




Assessment (optional) Repeat this analysis using the republican poll data. Filter for candidates (Trump, Cruz, Rubio, Carson, Bush) and for polls only after May 1, 2015. Color the lines by candidates.

Provide your code here

```
rep_polls %>%
  gather(key = candidate, value = percentage, Trump:Walker) %>%
  filter(candidate %in% c("Trump", "Cruz", "Rubio", "Carson", "Bush") &
    `End Date` >= "2015-05-01") %>%
  ggplot(aes(x=`End Date`, y = percentage, color = candidate)) +
  geom_line()
```

Warning: Removed 1 rows containing missing values (geom_path).



spread()

In contrast to *gathering* multiple columns into key-value pairs, we can *spread* a key-value pair across multiple columns.

The function `spread()` does just that. It transforms data from a *long* format into a *wide* format.

- The **key** is the name of the column in your data set that contains the values of the column headings that you are spreading across multiple columns
- The **value** is the name of the column that contains the values for the multiple columns

```
dem_polls_gathered <- dem_polls %>%
  gather(key = candidate, value = percentage,
         Clinton:Undecided)
dem_polls_gathered
```

```
## Source: local data frame [648 x 4]
##
##           Pollster   End Date candidate percentage
##           (chr)      (date)   (chr)      (int)
## 1 Emerson College Polling Society 2016-01-31 Clinton      51
## 2 Quinnipiac 2016-01-31 Clinton      46
## 3 Des Moines Register/Bloomberg/Selzer 2016-01-29 Clinton      45
## 4 Gravis Marketing/One America News 2016-01-27 Clinton      53
```



```
## 5          PPP (D-Progress Iowa) 2016-01-27  Clinton      48
## 6          NBC/WSJ/Marist 2016-01-26    Clinton      48
## 7          Monmouth University 2016-01-26    Clinton      47
## 8          ARG 2016-01-24      Clinton      45
## 9          Quinnipiac 2016-01-24      Clinton      45
## 10         Iowa State/WHO-HD 2016-01-22    Clinton      47
## ..          ...          ...          ...          ...
```

```
dem_polls_gathered %>%
  spread(key = candidate, value = percentage)
```

```
## Source: local data frame [81 x 10]
##
##      Pollster   End Date Biden Chafee Clinton Lessig O'Malley Sanders
##      (chr)      (date) (int)  (int)  (int)  (int)  (int)  (int)
## 1      ARG 2016-01-10    NA     NA     44     NA     3     47
## 2      ARG 2016-01-24    NA     NA     45     NA     3     48
## 3 CBS/YouGov 2015-09-10    10      1     33     NA     5     43
## 4 CBS/YouGov 2015-10-22    NA      1     46      0     3     43
## 5 CBS/YouGov 2015-11-19    NA     NA     50     NA     5     44
## 6 CBS/YouGov 2015-12-17    NA     NA     50     NA     4     45
## 7 CBS/YouGov 2016-01-21    NA     NA     46     NA     5     47
## 8      CNN 2014-09-10     15     NA     53     NA     2      5
## 9      CNN 2015-08-11     12      0     50     NA     1     31
## 10     CNN 2015-11-04     NA     NA     55     NA     3     37
## ..          ...          ...     ...     ...     ...     ...     ...
## Variables not shown: Undecided (int), Webb (int)
```

Other supporting functions in tidyr

- `separate()` = separate one column into multiple columns
- `unite()` = unite multiple columns into one

```
dem_polls_separate <- dem_polls %>%
  separate(col = `End Date`, into = c("y", "m", "d"))
dem_polls_separate
```

```
## Source: local data frame [81 x 12]
##
##      Pollster   y   m   d Clinton Sanders
##      (chr) (chr) (chr) (chr)  (int)  (int)
## 1 Emerson College Polling Society 2016 01 31     51     43
## 2 Quinnipiac 2016 01 31     46     49
## 3 Des Moines Register/Bloomberg/Selzer 2016 01 29     45     42
## 4 Gravis Marketing/One America News 2016 01 27     53     42
## 5 PPP (D-Progress Iowa) 2016 01 27     48     40
## 6 NBC/WSJ/Marist 2016 01 26     48     45
## 7 Monmouth University 2016 01 26     47     42
## 8 ARG 2016 01 24     45     48
## 9 Quinnipiac 2016 01 24     45     49
## 10 Iowa State/WHO-HD 2016 01 22     47     45
## ..          ...     ...     ...     ...     ...
```

```
## Variables not shown: O'Malley (int), Biden (int), Chafee (int), Lessig
##   (int), Webb (int), Undecided (int)
```

Assessment Use the `unite()` function to create a new column titled “end_date” that combines the columns `y`, `m` and `d` together into a single column separated by the “/” character.

```
## Provide your code here
```

```
dem_polls_separate %>%
  unite(col = end_date, y, m, d, sep = "/")
```

```
## Source: local data frame [81 x 10]
```

```
##
##               Pollster   end_date Clinton Sanders
##               (chr)      (chr)    (int)   (int)
## 1      Emerson College Polling Society 2016/01/31      51      43
## 2                Quinnipiac 2016/01/31      46      49
## 3 Des Moines Register/Bloomberg/Selzer 2016/01/29      45      42
## 4      Gravis Marketing/One America News 2016/01/27      53      42
## 5                PPP (D-Progress Iowa) 2016/01/27      48      40
## 6                NBC/WSJ/Marist 2016/01/26      48      45
## 7      Monmouth University 2016/01/26      47      42
## 8                ARG 2016/01/24      45      48
## 9                Quinnipiac 2016/01/24      45      49
## 10             Iowa State/WHO-HD 2016/01/22      47      45
## ..
## Variables not shown: O'Malley (int), Biden (int), Chafee (int), Lessig
##   (int), Webb (int), Undecided (int)
```

Cheatsheets

- [Data Wrangling with dplyr and tidyr from RStudio](#)